

# Customer Claim Prediction *at* Singlife Travel Insurance

Presented by: Sulthan Mahdi M. D.



# Table of contents

**01**

Business  
Problem

**02**

Data  
Understanding

**03**

Data Cleaning &  
Preprocessing

**04**

Analytics  
(Modeling)

**05**

Conclusion

**06**

Recommendation





01



# Business Problem





SingLife Travel Insurance, based in **Singapore**, is a **digital life insurance company** renowned for its innovative financial products. Specializing in the insurance sector, the company offers a range of digital insurance products, with a particular emphasis on **travel insurance**. Their coverage caters to both leisure and business travelers, reflecting their commitment to providing comprehensive and convenient insurance solutions.

# Problem and Goals

Despite SingLife's commitment to delivering insurance services, it has **encountered challenges** in efficiently **handling customer insurance claims**. To address this issue, the company is **seeking the expertise of a data scientist**. The data scientist's role would involve **analyzing** and **predicting** whether a customer is likely to file an insurance claim. This predictive analysis **aims to enhance the company's claim handling processes**, ensuring more effective and timely responses to customer needs.

Create a Machine Learning model that can be used by insurance companies to predict which customers will claim/not claim. Machine learning Models must be able to minimize losses to the smallest possible extent.

# Metrics Evaluation

– **FN:** The company estimates a loss of 800 SGD for cases where customers who actually make a claim (default) go undetected. This figure encompasses the loss from both the enrolled customers who file claims and the marketing costs incurred in acquiring new customers.

– **FP:** Meanwhile, the loss for situations where customers do not make a claim but are erroneously identified as claimants is 4000 SGD. This amount represents the average insurance coverage cost for these customers.

**We need to use fscore metrics with beta 1/5, this is because the precision value is 5x more important than the recall value**



# Data UnderStanding



02

# Data

<b>1. Agency</b>	Name of agency
<b>2. Agency Type</b>	Type of travel insurance agencies
<b>3. Distribution Channel</b>	Channel of travel insurance agencies
<b>4. Product Name</b>	Name of the travel insurance products
<b>5. Gender</b>	Gender of insured
<b>6. Duration</b>	Duration of travel
<b>7. Destination</b>	Destination of travel
<b>8. Net Sales</b>	Amount of sales of travel insurance policies
<b>9. Commission (in value)</b>	Commission received for travel insurance
<b>10. Age</b>	Age of insured
<b>11. Claim</b>	Claim status



# Data

	Agency	Agency Type	Distribution Channel	Product Name	Gender	Duration	Destination	Net Sales	Commision (in value)	Age	Claim
0	C2B	Airlines	Online	Annual Silver Plan	F	365	SINGAPORE	216.0	54.0	57	No
1	EPX	Travel Agency	Online	Cancellation Plan	NaN	4	MALAYSIA	10.0	0.0	33	No
2	JZI	Airlines	Online	Basic Plan	M	19	INDIA	22.0	7.7	26	No
3	EPX	Travel Agency	Online	2 way Comprehensive Plan	NaN	20	UNITED STATES	112.0	0.0	59	No
4	C2B	Airlines	Online	Bronze Plan	M	8	SINGAPORE	16.0	4.0	28	No

RangeIndex: 44328 entries, 0 to 44327

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Agency	44328 non-null	object
1	Agency Type	44328 non-null	object
2	Distribution Channel	44328 non-null	object
3	Product Name	44328 non-null	object
4	Gender	12681 non-null	object
5	Duration	44328 non-null	int64
6	Destination	44328 non-null	object
7	Net Sales	44328 non-null	float64
8	Commision (in value)	44328 non-null	float64
9	Age	44328 non-null	int64
10	Claim	44328 non-null	object

dtypes: float64(2), int64(2), object(7)

```
df.isna().sum()/len(df)*100
```

✓ 0.0s

Agency	0.000000
Agency Type	0.000000
Distribution Channel	0.000000
Product Name	0.000000
Gender	71.392799
Duration	0.000000
Destination	0.000000
Net Sales	0.000000
Commision (in value)	0.000000
Age	0.000000
Claim	0.000000

dtype: float64

```
df["Claim"].value_counts()
```

✓ 0.0s

Claim

No 38651


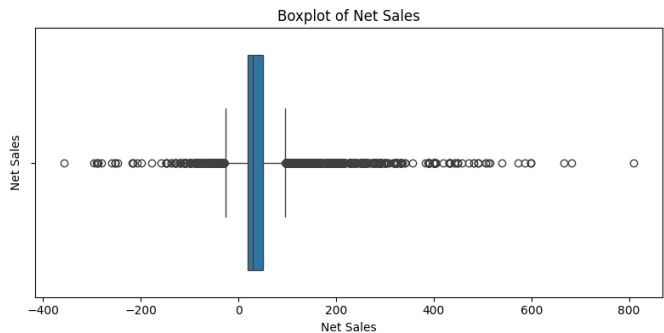
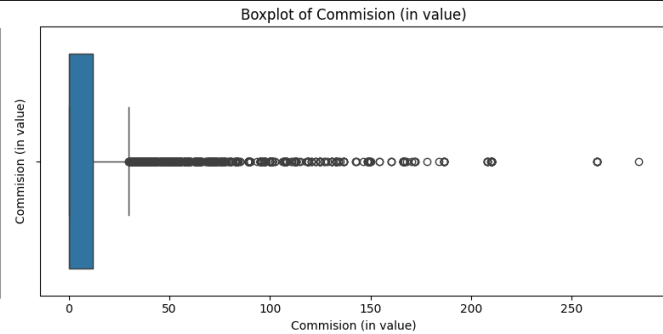
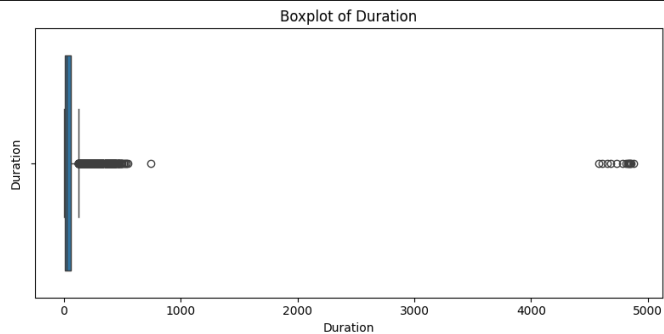

Yes 673

Name: count, dtype: int64



**03**

# **Data Cleaning & Preprocessing**


# Numerical Before Cleaning



# Numerical After Cleaning




# Categorical - Product Name



	Product_Name	Total_Claim	Cum_Percentage
0	Bronze Plan	159	25.52
1	Silver Plan	107	42.70
2	Annual Silver Plan	101	58.91
3	2 way Comprehensive Plan	94	74.00
4	Rental Vehicle Excess Insurance	54	82.66
5	Cancellation Plan	31	87.64
6	Basic Plan	16	90.21
7	Gold Plan	12	92.13

# Categorical - Destination



	Destination_Name	Total_Claim	Cum_Percentage
0	SINGAPORE	393	63.08
1	CHINA	29	67.74
2	AUSTRALIA	25	71.75
3	THAILAND	23	75.44
4	UNITED STATES	19	78.49
5	MALAYSIA	16	81.06
6	KOREA, REPUBLIC OF	12	82.99
7	UNITED KINGDOM	12	84.91
8	INDONESIA	10	86.52
9	JAPAN	10	88.12
10	VIET NAM	9	89.57
11	HONG KONG	8	90.85
12	PHILIPPINES	6	91.81
13	ITALY	6	92.78
14	TAIWAN, PROVINCE OF CHINA	5	93.58
15	FRANCE	5	94.38

# Column Transformer

```
transformer = ColumnTransformer([
    ("binary_encoding", BinaryEncoder(), ["Agency", "Product Name", "Destination"]),
    ("onehot_encoding", OneHotEncoder(drop="first"), ["Agency Type", "Distribution Channel"]),
    ("robust_scaling", RobustScaler(), ['Duration', 'Net Sales', 'Commision (in value)', 'Age'])
], remainder= "passthrough")
transformer
```

✓ 0.0s

Python

ColumnTransformer			
binary_encoding	onehot_encoding	robust_scaling	remainder
['Agency', 'Product Name', 'Destination']	['Agency Type', 'Distribution Channel']	['Duration', 'Net Sales', 'Commision (in value)', 'Age']	
BinaryEncoder	OneHotEncoder	RobustScaler	passthrough
BinaryEncoder()	OneHotEncoder(drop='first')	RobustScaler()	passthrough



# **Analytics (Modeling)**



**04**



# Cross Validation

```
f1per5_scorer = make_scorer(fbeta_score, beta=1/5)
```

```
# Cross Validation
```

```
models = [logreg, knn, tree, voting, stacking, bagging, rf, adaboost, gboost]
```

```
for i in models:
```

```
    pipe_model = Pipeline([
        ("preprocessing", transformer),
        ("modeling", i)
    ])
```

```
    model_cv = cross_val_score(
        estimator= pipe_model,
        X = X_train,
        y = y_train,
        cv = 5,
        scoring = f1per5_scorer,
        error_score = "raise",
        n_jobs = -1
    )
```

# Cross Validation Score

	algo	all_score	mean_score	std_score
0	LogisticRegression(random_state=0)	[0.7303, 0.8016, 0.6648, 0.7973, 0.7758]	0.753945	0.051268
4	StackingClassifier(estimators=[('clf1', Logist...	[0.7383, 0.8068, 0.6557, 0.7937, 0.7671]	0.752311	0.053740
3	VotingClassifier(estimators=[('clf1', Logistic...	[0.6935, 0.8016, 0.651, 0.7875, 0.7535]	0.737410	0.057083
8	GradientBoostingClassifier(random_state=0)	[0.6603, 0.7785, 0.6437, 0.7818, 0.7405]	0.720947	0.058388
7	AdaBoostClassifier(random_state=0)	[0.6758, 0.7531, 0.6478, 0.7403, 0.7241]	0.708248	0.039992
5	BaggingClassifier(estimator=KNeighborsClassifi...	[0.6765, 0.7698, 0.6305, 0.7362, 0.7255]	0.707670	0.048826
6	RandomForestClassifier(random_state=0)	[0.6683, 0.7863, 0.6384, 0.7403, 0.7049]	0.707639	0.052156
1	KNeighborsClassifier()	[0.6727, 0.7758, 0.6178, 0.7438, 0.7181]	0.705655	0.055358
2	DecisionTreeClassifier(random_state=0)	[0.6054, 0.6864, 0.6185, 0.6347, 0.6453]	0.638058	0.027728

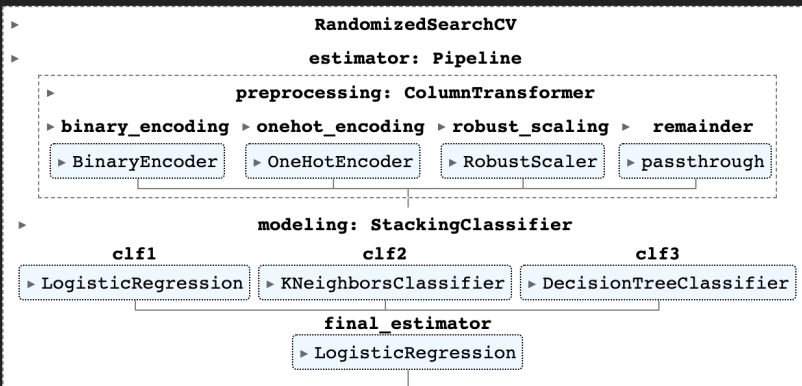
# Hyperparameter Tuning

```
# Hyperparameter tuning
hyperparam = {
    "modeling_clf1_C": np.logspace(5, -5, 11),
    "modeling_clf1_solver": ['lbfgs', 'liblinear'],
    "modeling_clf2_n_neighbors": range(1, 50, 2),
    "modeling_clf2_weights": ['uniform', 'distance'],
    "modeling_clf3_criterion": ['gini', 'entropy', 'log_loss'],
    "modeling_clf3_max_depth": range(2, 50, 1),
    "modeling_clf3_min_samples_split": range(2, 50),
    "modeling_clf3_min_samples_leaf": range(1, 25),
    "modeling_final_estimator": [logreg, knn, tree]
}

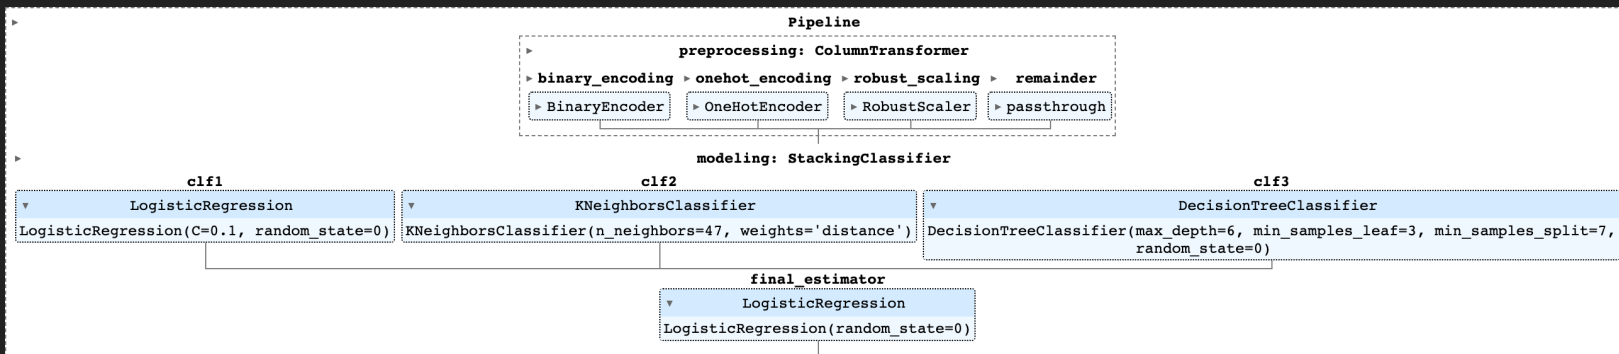
pipe_model = Pipeline([
    ("preprocessing", transformer),
    ("modeling", stacking)
])

randomsearch = RandomizedSearchCV(
    estimator=pipe_model,
    cv=5,
    n_jobs=-1,
    scoring=f1per5_scorer,
    param_distributions=hyperparam,
    random_state=0,
    n_iter=10000
)

randomsearch
```



# Hyperparameter Tuning Best Estimator



# Best Threshold

	threshold	f1/5 train	f1/5 test
16	0.76	0.974625	0.856654
12	0.72	0.978815	0.853682
7	0.67	0.965551	0.853458
5	0.65	0.965119	0.852700
2	0.62	0.965354	0.850701

0.7948411750656795 No treatment

0.802469135802469 Parameter Tuning


0.8566538296961916 Parameter Tuning + Optimized Threshold



05


# Conclusion

# Confusion Matrix




**Before Tuning:**

– FP:  
 $32 \times 4000 \text{ SGD} = 128000 \text{ SGD}$   
– FN:  
 $59 \times 800 \text{ SGD} = 47200 \text{ SGD}$   
**Total loss: 175200 SGD**



**After Tuning:**

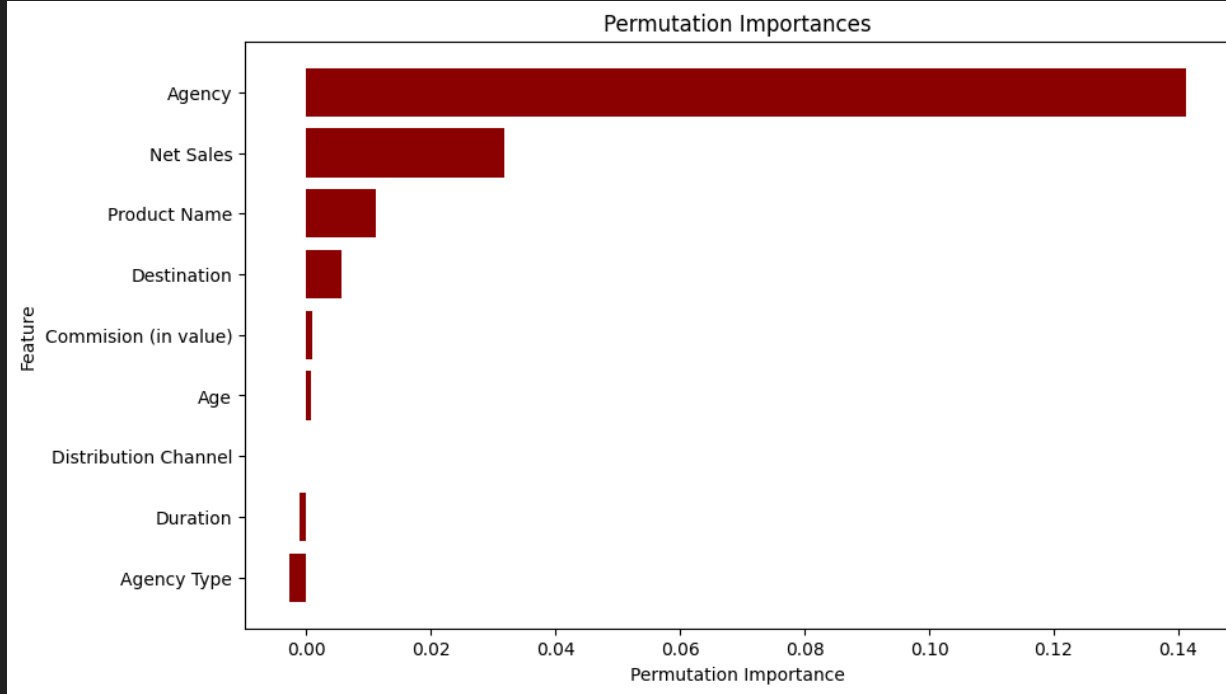
– FP:  
 $31 \times 4000 \text{ SGD} = 124000 \text{ SGD}$   
– FN:  
 $57 \times 800 \text{ SGD} = 45600 \text{ SGD}$   
**Total loss: 169600 SGD**



**After Tuning + Threshold:**

– FP:  
 $9 \times 4000 \text{ SGD} = 36000 \text{ SGD}$   
– FN:  
 $110 \times 800 \text{ SGD} = 88000 \text{ SGD}$   
**Total loss: 124000 SGD**

# Feature Importance







# Recommendation



06



# Recommendations

Recommendations to improve model performance:

- Perform Hyperparameter tuning with GridSearchCV to find the optimal parameter combination in the model.
- Add data to the Claim target in the positive class (Claim = 1).
- Try to handle imbalance using other methods besides Random Under Sampling.
- Try using other metrics such as a combination of Recall and FPR.



# Thank You!